






**Prefill: 100 tokens, Decoding: 100 tokens
(200 reqs/config)**




**Prefill: 100 tokens, Decoding: 200 tokens
(200 reqs/config)**




**Prefill: 200 tokens, Decoding: 100 tokens
(200 reqs/config)**



**Prefill: 200 tokens, Decoding: 200 tokens
(200 reqs/config)**



**Prefill: 400 tokens, Decoding: 100 tokens
(200 reqs/config)**



**Prefill: 50 tokens, Decoding: 200 tokens
(200 reqs/config)**

